

# Aircraft Analysis from Aviation Accident Data.

## Introduction

Our company is entering the aviation industry, and we need to make smart choices to keep risks low. This project uses Aviation Accident data from the U.S. National Transportation Safety Board (NTSB) to help us make good buying decisions.

I will do this by carefully looking at, cleaning, and studying the data. The information I will find will give clear advice to the head of our new aviation department. This advice will help pick aircraft that are safer, allowing our company to start this new business with a stronger and more secure beginning.

## Data Understanding

This project relies on Aviation Accident data from [Kaggle](#), originally from the U.S. National Transportation Safety Board (NTSB). For full project context and key questions, refer to the [README](#).

## 1. Data Exploration

With the project clearly explained, i will now load and explore the dataset to understand its structure, size, and contents.

```
# Import relevant libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

plt.style.available
plt.style.use('seaborn-darkgrid')
```

Next is Loading the dataset and previewing the first 5 and last 5 records

```
# Load the aviation dataset
df = pd.read_csv('Data/AviationData.csv', encoding = 'latin1',
low_memory = False)

# Display the first 5 rows
df.head()
```

	Event.Id	Investigation.Type	Accident.Number	Event.Date	\
0	20001218X45444	Accident	SEA87LA080	1948-10-24	
1	20001218X45447	Accident	LAX94LA336	1962-07-19	

2	20061025X01555	Accident	NYC07LA005	1974-08-30
3	20001218X45448	Accident	LAX96LA321	1977-06-19
4	20041105X01764	Accident	CHI79FA064	1979-08-02

	Location	Country	Latitude	Longitude	Airport.Code
\					
0	MOOSE CREEK, ID	United States	NaN	NaN	NaN
1	BRIDGEPORT, CA	United States	NaN	NaN	NaN
2	Saltville, VA	United States	36.922223	-81.878056	NaN
3	EUREKA, CA	United States	NaN	NaN	NaN
4	Canton, OH	United States	NaN	NaN	NaN

	Airport.Name	...	Purpose.of.flight	Air.carrier	Total.Fatal.Injuries
\					
0	NaN	...	Personal	NaN	2.0
1	NaN	...	Personal	NaN	4.0
2	NaN	...	Personal	NaN	3.0
3	NaN	...	Personal	NaN	2.0
4	NaN	...	Personal	NaN	1.0

	Total.Serious.Injuries	Total.Minor.Injuries	Total.Uninjured	\
0	0.0	0.0	0.0	
1	0.0	0.0	0.0	
2	NaN	NaN	NaN	
3	0.0	0.0	0.0	
4	2.0	NaN	0.0	

	Weather.Condition	Broad.phase.of.flight	Report.Status
Publication.Date			
0	UNK	Cruise	Probable Cause
NaN			
1	UNK	Unknown	Probable Cause
09-1996			
2	IMC	Cruise	Probable Cause
02-2007			
3	IMC	Cruise	Probable Cause
09-2000			
4	VMC	Approach	Probable Cause
04-1980			

[5 rows x 31 columns]

*# Check the summary of the dataset to see the number of rows, columns, data types, and any missing values.*

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 88889 entries, 0 to 88888
```

```
Data columns (total 31 columns):
```

#	Column	Non-Null Count	Dtype
0	Event.Id	88889 non-null	object
1	Investigation.Type	88889 non-null	object
2	Accident.Number	88889 non-null	object
3	Event.Date	88889 non-null	object
4	Location	88837 non-null	object
5	Country	88663 non-null	object
6	Latitude	34382 non-null	object
7	Longitude	34373 non-null	object
8	Airport.Code	50249 non-null	object
9	Airport.Name	52790 non-null	object
10	Injury.Severity	87889 non-null	object
11	Aircraft.damage	85695 non-null	object
12	Aircraft.Category	32287 non-null	object
13	Registration.Number	87572 non-null	object
14	Make	88826 non-null	object
15	Model	88797 non-null	object
16	Amateur.Built	88787 non-null	object
17	Number.of.Engines	82805 non-null	float64
18	Engine.Type	81812 non-null	object
19	FAR.Description	32023 non-null	object
20	Schedule	12582 non-null	object
21	Purpose.of.flight	82697 non-null	object
22	Air.carrier	16648 non-null	object
23	Total.Fatal.Injuries	77488 non-null	float64
24	Total.Serious.Injuries	76379 non-null	float64
25	Total.Minor.Injuries	76956 non-null	float64
26	Total.Uninjured	82977 non-null	float64
27	Weather.Condition	84397 non-null	object
28	Broad.phase.of.flight	61724 non-null	object
29	Report.Status	82508 non-null	object
30	Publication.Date	75118 non-null	object

```
dtypes: float64(5), object(26)
```

```
memory usage: 21.0+ MB
```

*# Generate descriptive statistics to get an overview of the distributions and general characteristics of the data in the cell*

```
df.describe()
```

	Number.of.Engines	Total.Fatal.Injuries	Total.Serious.Injuries
\			
count	82805.000000	77488.000000	76379.000000
mean	1.146585	0.647855	0.279881
std	0.446510	5.485960	1.544084
min	0.000000	0.000000	0.000000
25%	1.000000	0.000000	0.000000
50%	1.000000	0.000000	0.000000
75%	1.000000	0.000000	0.000000
max	8.000000	349.000000	161.000000

	Total.Minor.Injuries	Total.Uninjured
count	76956.000000	82977.000000
mean	0.357061	5.325440
std	2.235625	27.913634
min	0.000000	0.000000
25%	0.000000	0.000000
50%	0.000000	1.000000
75%	0.000000	2.000000
max	380.000000	699.000000

```
# Explore the unique values in each column of our dataset
print("\nNumber of unique values in each column:")
df.nunique()
```

Number of unique values in each column:

Event.Id	87951
Investigation.Type	2
Accident.Number	88863
Event.Date	14782
Location	27758
Country	219
Latitude	25589
Longitude	27154
Airport.Code	10375
Airport.Name	24871
Injury.Severity	109
Aircraft.damage	4
Aircraft.Category	15
Registration.Number	79105
Make	8237

Model	12318
Amateur.Built	2
Number.of.Engines	7
Engine.Type	13
FAR.Description	31
Schedule	3
Purpose.of.flight	26
Air.carrier	13590
Total.Fatal.Injuries	125
Total.Serious.Injuries	50
Total.Minor.Injuries	57
Total.Uninjured	379
Weather.Condition	4
Broad.phase.of.flight	12
Report.Status	17075
Publication.Date	2924
dtype: int64	

*# Check the datatypes in each column*  
df.dtypes

Event.Id	object
Investigation.Type	object
Accident.Number	object
Event.Date	object
Location	object
Country	object
Latitude	object
Longitude	object
Airport.Code	object
Airport.Name	object
Injury.Severity	object
Aircraft.damage	object
Aircraft.Category	object
Registration.Number	object
Make	object
Model	object
Amateur.Built	object
Number.of.Engines	float64
Engine.Type	object
FAR.Description	object
Schedule	object
Purpose.of.flight	object
Air.carrier	object
Total.Fatal.Injuries	float64
Total.Serious.Injuries	float64
Total.Minor.Injuries	float64
Total.Uninjured	float64
Weather.Condition	object
Broad.phase.of.flight	object

```
Report.Status      object
Publication.Date    object
dtype: object
```

```
# Check for missing values
```

```
df.isnull().sum()
```

```
# Calculate the percentage of missing values per column
```

```
null_percent = (df.isnull().sum() / len(df)) * 100
```

```
# Sort the columns from highest percentage of missing values to lowest.
```

```
null_percent.sort_values(ascending = False)
```

```
Schedule      85.845268
Air.carrier    81.271023
FAR.Description 63.974170
Aircraft.Category 63.677170
Longitude      61.330423
Latitude       61.320298
Airport.Code   43.469946
Airport.Name   40.611324
Broad.phase.of.flight 30.560587
Publication.Date 15.492356
Total.Serious.Injuries 14.073732
Total.Minor.Injuries 13.424608
Total.Fatal.Injuries 12.826109
Engine.Type    7.961615
Report.Status  7.178616
Purpose.of.flight 6.965991
Number.ofEngines 6.844491
Total.Uninjured 6.650992
Weather.Condition 5.053494
Aircraft.damage 3.593246
Registration.Number 1.481623
Injury.Severity 1.124999
Country        0.254250
Amateur.Built  0.114750
Model          0.103500
Make           0.070875
Location       0.058500
Event.Date     0.000000
Accident.Number 0.000000
Investigation.Type 0.000000
Event.Id       0.000000
dtype: float64
```

```
# Check for duplicates in the dataset
```

```
df.duplicated().sum()
```

## Conclusion on Data Exploration

The dataset contains a mix of categorical and numerical features related to aviation accidents. No duplicate records were found, but several columns contain missing values and will require cleaning. I will focus on selecting relevant columns for deeper analysis, with the goal of identifying patterns and risk factors that can support safer aircraft acquisition decisions.

## 2. Data Cleaning

Data cleaning is a crucial step to ensure the quality and reliability of the analysis. In this section i will handle missing values and correct inconsistent data formats. Clean data will help produce accurate insights that support better decision-making.

First is to standardize Column Names since some use dot like, `Event.Id`, `Total.Fatal.Injuries`. Let's clean them.

```
# Let's check the column names here
df.columns

Index(['Event.Id', 'Investigation.Type', 'Accident.Number',
      'Event.Date',
      'Location', 'Country', 'Latitude', 'Longitude', 'Airport.Code',
      'Airport.Name', 'Injury.Severity', 'Aircraft.damage',
      'Aircraft.Category', 'Registration.Number', 'Make', 'Model',
      'Amateur.Built', 'Number.of.Engines', 'Engine.Type',
      'FAR.Description',
      'Schedule', 'Purpose.of.flight', 'Air.carrier',
      'Total.Fatal.Injuries',
      'Total.Serious.Injuries', 'Total.Minor.Injuries',
      'Total.Uninjured',
      'Weather.Condition', 'Broad.phase.of.flight', 'Report.Status',
      'Publication.Date'],
      dtype='object')

# Cleaning now column names
df.columns = df.columns.str.capitalize().str.replace('.', '_', regex =
False)
df.columns

Index(['Event_id', 'Investigation_type', 'Accident_number',
      'Event_date',
      'Location', 'Country', 'Latitude', 'Longitude', 'Airport_code',
      'Airport_name', 'Injury_severity', 'Aircraft_damage',
      'Aircraft_category', 'Registration_number', 'Make', 'Model',
      'Amateur_built', 'Number_of_engines', 'Engine_type',
      'Far_description',
      'Schedule', 'Purpose_of_flight', 'Air_carrier',
      'Total_fatal_injuries',
```

```

        'Total_serious_injuries', 'Total_minor_injuries',
'Total_uninjured',
        'Weather_condition', 'Broad_phase_of_flight', 'Report_status',
        'Publication_date'],
        dtype='object')

```

```

# Check each column dtypes
df.dtypes

```

```

Event_id                object
Investigation_type      object
Accident_number         object
Event_date              object
Location                object
Country                 object
Latitude                object
Longitude               object
Airport_code            object
Airport_name            object
Injury_severity         object
Aircraft_damage         object
Aircraft_category       object
Registration_number     object
Make                    object
Model                   object
Amateur_built           object
Number_of_engines       float64
Engine_type             object
Far_description          object
Schedule                object
Purpose_of_flight       object
Air_carrier             object
Total_fatal_injuries    float64
Total_serious_injuries  float64
Total_minor_injuries    float64
Total_uninjured         float64
Weather_condition       object
Broad_phase_of_flight   object
Report_status           object
Publication_date        object
dtype: object

```

Examining the columns, we see that the `Event_date` column is currently stored as an object data type. We convert it into a proper datetime format using the `pd.to_datetime()` function to enable time-based analysis.

```

# Convert Event_date to datetime
df['Event_date'] = pd.to_datetime(df['Event_date'], errors = 'coerce')

```



```

# Check for dates not converted
print(f"The number not converted is :
{df['Event_date'].isnull().sum()}")

# Extract Year to help in analysis
df['Year'] = df['Event_date'].dt.year

# Extract Month also
df['Month'] = df['Event_date'].dt.month

print(df[['Event_date', 'Year', 'Month']].head())

```

```

The number not converted is : 0
  Event_date  Year  Month
0 1948-10-24  1948     10
1 1962-07-19  1962      7
2 1974-08-30  1974      8
3 1977-06-19  1977      6
4 1979-08-02  1979      8

```

Since injury columns have missing values, they likely indicate no injuries occurred or unreported incidents and they are `float64` and might have NaNs. I will fill NaNs with 0 and convert injury columns to integers.

```

# Columns to clean and convert to integers
injury_cols = ['Total_fatal_injuries', 'Total_serious_injuries',
'Total_minor_injuries', 'Total_uninjured']

for col in injury_cols:
    # convert to numeric
    df[col] = pd.to_numeric(df[col], errors = 'coerce')
    # Fill missing values with 0
    df[col].fillna(0, inplace = True)
    # Convert to integer
    df[col] = df[col].astype(int)

print("\nInjury columns after conversion")
print(df[injury_cols].info())

```

```

Injury columns after conversion
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 88889 entries, 0 to 88888
Data columns (total 4 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Total_fatal_injuries                  88889 non-null  int32
1   Total_serious_injuries                88889 non-null  int32
2   Total_minor_injuries                  88889 non-null  int32

```

```
3    Total_uninjured      88889 non-null  int32
dtypes: int32(4)
memory usage: 1.4 MB
None
```

```
# Check missing values in each column again
df.isnull().sum()
```

```
Event_id      0
Investigation_type  0
Accident_number  0
Event_date    0
Location      52
Country       226
Latitude      54507
Longitude     54516
Airport_code  38640
Airport_name  36099
Injury_severity  1000
Aircraft_damage  3194
Aircraft_category  56602
Registration_number  1317
Make          63
Model         92
Amateur_built  102
Number_of_engines  6084
Engine_type    7077
Far_description  56866
Schedule       76307
Purpose_of_flight  6192
Air_carrier    72241
Total_fatal_injuries  0
Total_serious_injuries  0
Total_minor_injuries  0
Total_uninjured  0
Weather_condition  4492
Broad_phase_of_flight  27165
Report_status   6381
Publication_date  13771
Year            0
Month           0
dtype: int64
```

We now handle the `Make` and `Model` columns since they define our aircraft type hence there will be need to combine them. Also potential missing values will be dealt with.

```
# To minimize errors during concatenation i convert the columns to string
df['Make'] = df['Make'].astype(str)
```

```

df['Model'] = df['Model'].astype(str)
# Combine 'Make' and 'Model' into a new 'Aircraft_Type' column
df['Aircraft_Type'] = df['Make'] + ' ' + df['Model']
# Replace 'nan' string values that result from missing data in
original Make/Model
df['Aircraft_Type'] = df['Aircraft_Type'].replace('nan nan', np.nan)
# Drop rows where 'Aircraft_Type' is still NaN
df.dropna(subset = ['Aircraft_Type'], inplace = True)

print(f"DataFrame shape after dropping missing Aircraft_Type:
{df.shape}")

# Check unique values for 'Aircraft_Type'
print(f"Number of unique values: {df['Aircraft_Type'].nunique()}")
print(df['Aircraft_Type'].value_counts().head(10))

DataFrame shape after dropping missing Aircraft_Type: (88846, 34)
Number of unique values: 20182
Cessna 152          2168
Cessna 172          1254
Cessna 172N          996
Piper PA-28-140      812
Cessna 150           716
Cessna 172M          667
Cessna 172P          597
Piper PA-18          539
Cessna 150M          539
Piper PA-28-161      502
Name: Aircraft_Type, dtype: int64

```

Next we are going to clean `Weather_Condition` and `Broad_phase_of_flights` since they are categorical and important. We'll check the unique values and consider standardization.

```

# Check 'Weather_Condition'
df['Weather_condition'].value_counts(dropna=False)

df['Weather_condition'] = df['Weather_condition'].replace({'UNK':
'Unknown', 'Unk': 'Unknown'})
df['Weather_condition'].fillna('Unknown', inplace=True)

print(df['Weather_condition'].value_counts(dropna=False))

VMC          77295
IMC           5976
Unknown       5575
Name: Weather_condition, dtype: int64

# Check 'Broad_phase_of_flight'
df['Broad_phase_of_flight'].value_counts(dropna=False)

```

```

df['Broad_phase_of_flight'] =
df['Broad_phase_of_flight'].replace({'Unk': 'Unknown'})
df['Broad_phase_of_flight'].fillna('Unknown', inplace=True) # Fill any
actual NaNs

print(df['Broad_phase_of_flight'].value_counts(dropna=False))

```

Unknown	27670
Landing	15428
Takeoff	12493
Cruise	10269
Maneuvering	8144
Approach	6546
Climb	2034
Taxi	1958
Descent	1887
Go-around	1353
Standing	945
Other	119

```

Name: Broad_phase_of_flight, dtype: int64

```

Now i will drop the original `Make` and `Model` columns as they are now combined into `Aircraft_Type` column. Also considering dropping other columns that are not directly used or have too many NaNs for this analysis.

```

# List of columns to drop.
columns_drop = [
    'Make', 'Model', 'Accident_number', 'Investigation_type',
    'Event_id', 'Latitude', 'Longitude', 'Airport_code', 'Airport_name',
    'Publication_date', 'Report_status', 'Far_description',
    'Aircraft_damage', 'Injury_severity', 'Aircraft_category',
    'Air_carrier',
    'Engine_type', 'Schedule', 'Registration_number', 'Amateur_built']
# drop the columns
df.drop(columns=columns_drop, inplace=True, errors='ignore')

print(f"Dataframe after dropped columns : {df.shape}")

Dataframe after dropped columns : (88846, 14)

# Check for missing values again
df.isnull().sum()

```

Event_date	0
Location	52
Country	225
Number_of_engines	6043
Purpose_of_flight	6153
Total_fatal_injuries	0

Total_serious_injuries	0
Total_minor_injuries	0
Total_uninjured	0
Weather_condition	0
Broad_phase_of_flight	0
Year	0
Month	0
Aircraft_Type	0
dtype: int64	

Looking at the missing values above the `Number_of_engines` column still has missing values. Knowing if single-engine aircraft are riskier than multi-engine ones could be a key insight for your company. So i will proceed to clean the column.

```
# Clean the column
df['Number_of_engines'].value_counts(dropna=False)

# For number of engines, mode makes sense as 1.0 and 2.0 are frequent
mode_engines = df['Number_of_engines'].mode()[0] # Gets the most frequent value
# Convert to int after filling NaNs
df['Number_of_engines'].fillna(mode_engines, inplace=True)
df['Number_of_engines'] = df['Number_of_engines'].astype(int)

print(df['Number_of_engines'].value_counts(dropna=False))
```

1	75624
2	11078
0	1226
3	483
4	431
8	3
6	1

Name: Number\_of\_engines, dtype: int64

Next we look at the `Purpose_of_flight` column. It has several missing values but it can provide insights if commercial passenger flights have different risk than personal flights.

```
# Clean Purpose_of_flight
df['Purpose_of_flight'].value_counts(dropna = False)
# Fill with unknown
df['Purpose_of_flight'].fillna('Unknown', inplace=True)

print("\n After Cleaning 'Purpose_of_flight' ")
print(df['Purpose_of_flight'].value_counts(dropna=False))
```

After Cleaning 'Purpose_of_flight'	
Personal	49446

Unknown	12953
Instructional	10601
Aerial Application	4712
Business	4018
Positioning	1646
Other Work Use	1264
Ferry	812
Aerial Observation	794
Public Aircraft	720
Executive/corporate	553
Flight Test	405
Skydiving	182
External Load	123
Public Aircraft - Federal	105
Banner Tow	101
Air Race show	99
Public Aircraft - Local	74
Public Aircraft - State	64
Air Race/show	59
Glider Tow	53
Firefighting	40
Air Drop	11
ASHO	6
PUBS	4
PUBL	1

Name: Purpose\_of\_flight, dtype: int64

*# Lets run the missing value check again to see columns that have issues*

```
df.isnull().sum()
```

Event_date	0
Location	52
Country	225
Number_of_engines	0
Purpose_of_flight	0
Total_fatal_injuries	0
Total_serious_injuries	0
Total_minor_injuries	0
Total_uninjured	0
Weather_condition	0
Broad_phase_of_flight	0
Year	0
Month	0
Aircraft_Type	0

dtype: int64

Lastly the **Location** and **Country** columns still has missing value. I will consider dropping the rows with missing value to ensure I'm working with records where the location and country are definitively known, which will be important for geographical filtering.

```

# Lets check the value_counts in the two columns
df[['Location', 'Country']].isnull().sum()

# Drop rows where 'Location' or 'Country' is missing
df.dropna(subset=['Location', 'Country'], inplace=True)

# Confirm if the missing values are still there
print(df[['Location', 'Country']].isnull().sum())

Location      0
Country       0
dtype: int64

#Final check on missing values
df.isnull().sum()

Event_date      0
Location        0
Country         0
Number_of_engines  0
Purpose_of_flight  0
Total_fatal_injuries  0
Total_serious_injuries  0
Total_minor_injuries  0
Total_uninjured  0
Weather_condition  0
Broad_phase_of_flight  0
Year            0
Month           0
Aircraft_Type   0
dtype: int64

# Save the modified DataFrame to a new Excel file
df.to_excel("clean_data.xlsx", index=False)

```

## Conclusion

The initial **Aviation Accident** dataset has been thoroughly cleaned and prepared. Key steps included standardizing column names, converting data types e.g dates to datetime, injuries columns to integers and creating a combined **Aircraft\_Type** column. Importantly, all missing values in critical columns have been addressed, resulting in a dataset of 88,570 entries across 14 relevant columns. This cleaned data is now fully ready for Exploratory Data Analysis to identify low-risk aircraft.

## Exploratory Data Analysis

Having cleaned the dataset, the next phase is Exploratory Data Analysis (EDA). EDA is an essential step in any data project, acting as a detective phase where we investigate the dataset's main characteristics and uncover patterns often with visual methods.

For this project, EDA will enable us to visually and statistically explore accident frequencies, fatality counts, and the influence of various factors like aircraft type, weather, and flight phase to directly address our company's goal of identifying the lowest-risk aircraft for acquisition.

## Overall Accident Trends

We begin by looking at the number of accidents caused by aircraft types with a spacing of 10 year interval.

```
# Create column for 10-year intervals
df['Years_10'] = (df['Year'] // 10) * 10

# Find top 10 aircraft counts then groupby and pivot
top_10_aircraft = df['Aircraft_Type'].value_counts().head(10).index
df_top = df[df['Aircraft_Type'].isin(top_10_aircraft)].copy()
accidents_top_10 = df_top.groupby(['Years_10',
'Aircraft_Type']).size().reset_index(name = 'Accident_Count')
pivot_top_10 = accidents_top_10.pivot(index='Years_10', columns =
'Aircraft_Type', values = 'Accident_Count').fillna(0)

# Find Low-accident aircraft counts then groupby and pivot Low-
accident aircraft
aircraft_counts = df['Aircraft_Type'].value_counts()
low_aircraft = aircraft_counts[(aircraft_counts >= 5) &
(aircraft_counts <= 15)].head(10).index
df_low = df[df['Aircraft_Type'].isin(low_aircraft)]
accidents_low = df_low.groupby(['Years_10',
'Aircraft_Type']).size().reset_index(name = 'Accident_Count')
pivot_low = accidents_low.pivot(index = 'Years_10', columns =
'Aircraft_Type', values = 'Accident_Count').fillna(0)

# Plotting
fig, axes = plt.subplots(1, 2, figsize = (16, 8), sharey = True)


# Plotting the top 10
pivot_top_10.plot(ax = axes[0], kind = 'line', marker = 'o')
axes[0].set_title("Top 10 Aircraft Types Accident Trends by 10 Years")
axes[0].set_xlabel("Years")
axes[0].set_ylabel("Number of Accidents")
axes[0].grid(True, linestyle = '--', alpha = 0.5)

# Plot low-accident
pivot_low.plot(ax = axes[1], kind = 'line', marker = 'o')
axes[1].set_title("Low-Accident Aircraft Types Trends by 10 Years")
axes[1].set_xlabel("Years")
axes[1].grid(True, linestyle = '--', alpha = 0.8)

# Layout and show
plt.tight_layout()
```



```
plt.savefig('Images/Trends.png', dpi=300, bbox_inches='tight') # saves
visual to image folder of my project
plt.show()
```



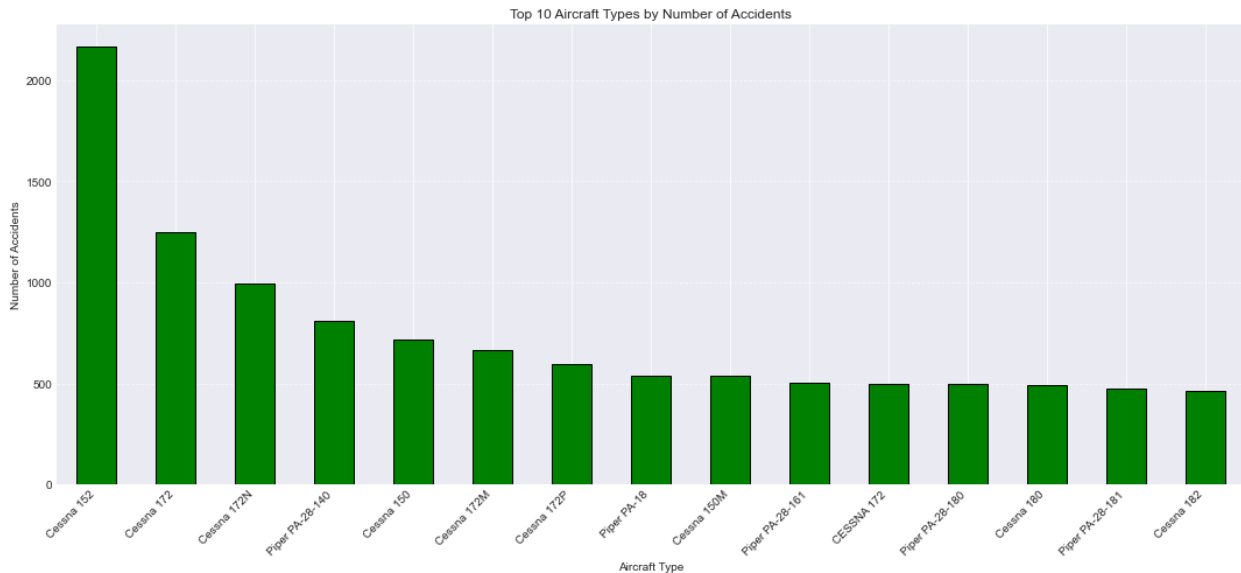
The accident trend analysis reveals that aircraft like Cessna 152, Cessna 172, and Piper PA-28-140 have high accident counts. In contrast, aircraft such as Piper PA-34, Taylorcraft DCO-65, and Boeing 737-222 show flat, low accident trends.

Aircraft types with high number of accidents

A bar chart of aircraft type vs number of accidents will help directly answer: Which aircraft types have been involved in the most accidents overall?

```
# Count total accidents per aircraft type
aircraft_accidents = df['Aircraft_Type'].value_counts().head(15)

# Plot
plt.figure(figsize = (15, 7))
aircraft_accidents.plot(kind = 'bar', color = 'green', edgecolor =
'black')
plt.title('Top 10 Aircraft Types by Number of Accidents')
plt.xlabel('Aircraft Type')
plt.ylabel('Number of Accidents')
plt.xticks(rotation = 45, ha = 'right')
plt.grid(axis = 'y', linestyle = '--', alpha = 0.7)
plt.tight_layout()
plt.savefig('Images/Accidents.png', dpi=300, bbox_inches='tight') #
saves visual to image folder of my project
plt.show()
```



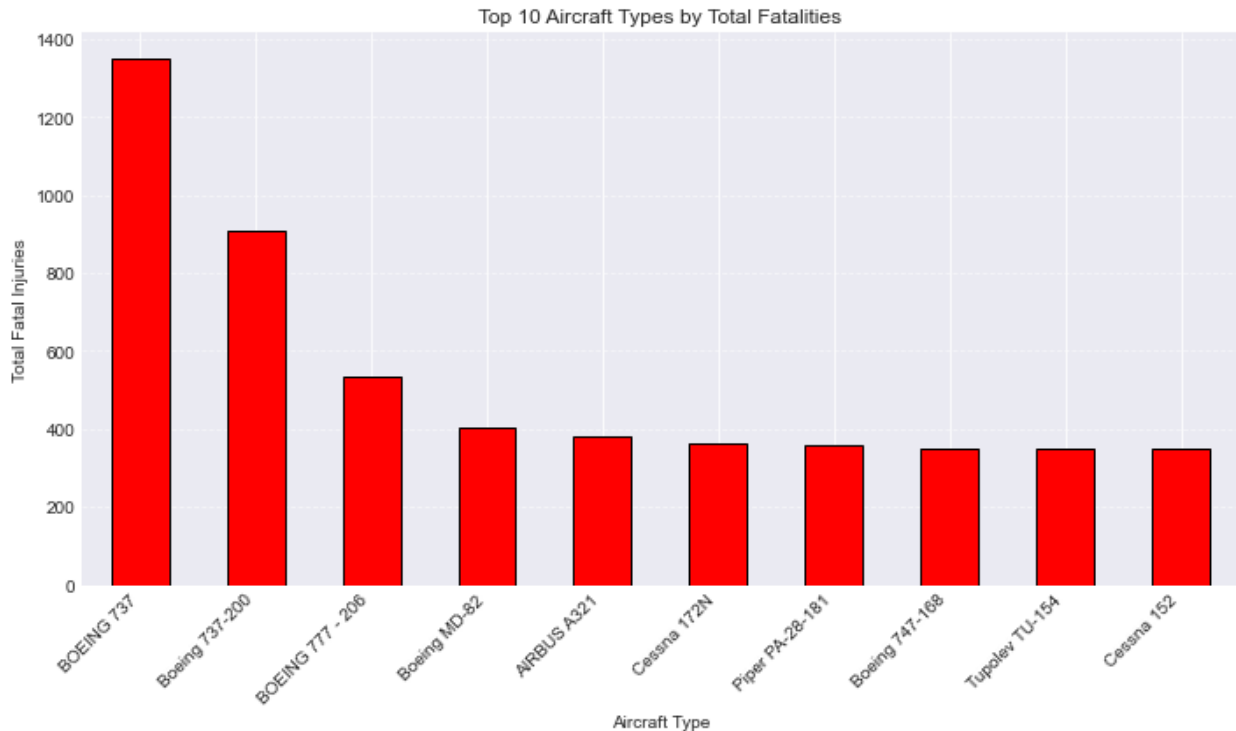
This chart highlights which aircraft have been most frequently involved in accidents. The Cessna 152 and Cessna 172 top the list.

## Fatal Injuries Across Aircraft Types

This visualization will use a barchart to compare the distribution of fatal injuries by aircraft types. The chart shows which aircraft types are deadliest based on total fatalities, even if they didn't crash often.

```
# Group by aircraft type and sum fatalities
fatal_df = df.groupby('Aircraft_Type')
['Total_fatal_injuries'].sum().sort_values(ascending = False).head(10)

# Plot
plt.figure(figsize = (10, 6))
fatal_df.plot(kind = 'bar', color = 'red', edgecolor = 'black')
plt.title("Top 10 Aircraft Types by Total Fatalities")
plt.xlabel("Aircraft Type")
plt.ylabel("Total Fatal Injuries")
plt.xticks(rotation = 45, ha = 'right')
plt.grid(axis = 'y', linestyle = '--', alpha = 0.6)
plt.tight_layout()
plt.savefig('Images/Fatalities.png', dpi=300, bbox_inches='tight') #
saves visual to image folder of my project
plt.show()
```



The analysis showed that the Boeing 737 and Boeing 737-200 accounted for the highest number of fatalities across all aircraft types in the dataset.

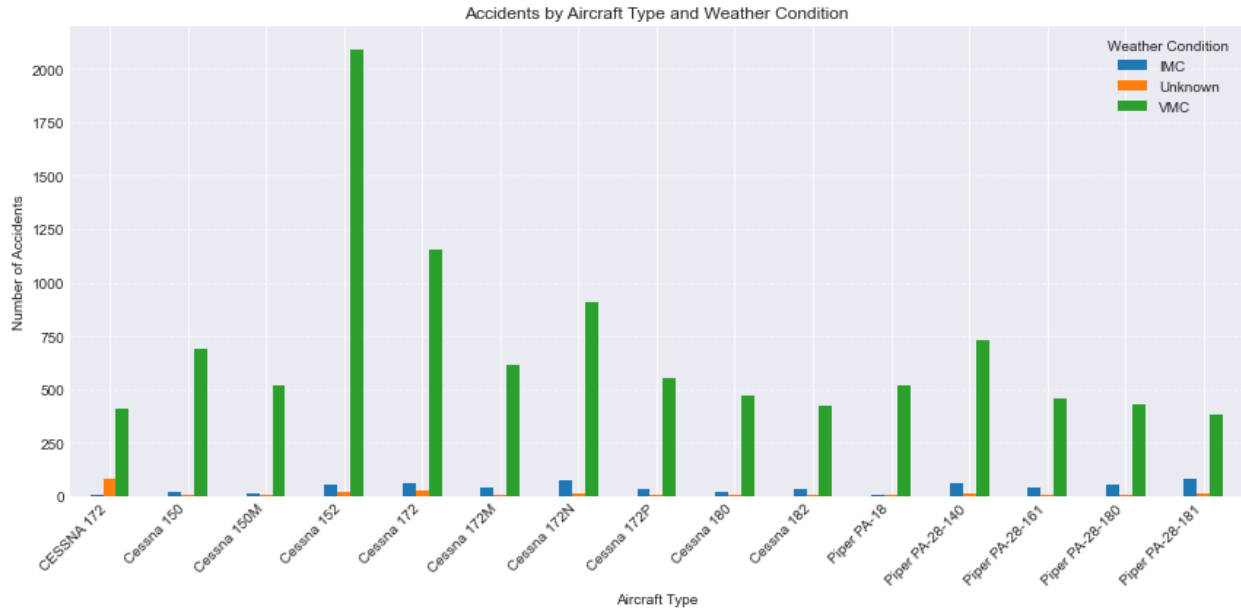
## Impact of Weather

How an external factor weather, affects the accident data. We will create a visualization to show the distribution of accidents by `Weather_condition`.

```
# Get top 15 aircrafts
top_15_aircraft = df['Aircraft_Type'].value_counts().head(15).index
df_top = df[df['Aircraft_Type'].isin(top_15_aircraft)]

# Group and pivot
weather_by_aircraft = df_top.groupby(['Aircraft_Type',
'Weather_condition']).size().unstack(fill_value = 0)

# Plot
weather_by_aircraft.plot(kind = 'bar', figsize = (12, 6))
plt.title('Accidents by Aircraft Type and Weather Condition')
plt.xlabel('Aircraft Type')
plt.ylabel('Number of Accidents')
plt.xticks(rotation = 45, ha = 'right')
plt.legend(title = 'Weather Condition')
plt.tight_layout()
plt.grid(axis = 'y', linestyle = '--', alpha = 0.6)
plt.savefig('Images/Weather.png', dpi=300, bbox_inches='tight') #
saves visual to image folder of my project
plt.show()
```



The chart shows that Cessna 152 and Cessna 172 have the highest number of accidents under (VMC).

## Accident Distribution by Broad Phase of Flight

Shows which phases of flight like Landing, Takeoff, Cruise are most commonly associated with accidents for each aircraft.

```
# Clean and standardize the phase column
df['Phase_broad'] =
df['Broad_phase_of_flight'].str.upper().str.strip()

# Get the top 4 most common phases
top_phases = df['Phase_broad'].value_counts().head(4).index

# Prepare figure layout
fig, axes = plt.subplots(2, 2, figsize=(16, 10))
axes = axes.flatten() # Make it easy to iterate over

# Loop over each phase and create individual bar plot
for i, phase in enumerate(top_phases):
    # Filter to only rows for this phase
    df_phase = df[df['Phase_broad'] == phase]

    # Group by aircraft type and count accidents
    top5_aircraft = (df_phase['Aircraft_Type'].value_counts().head(5))


    # Plot
    sns.barplot(x = top5_aircraft.values, y = top5_aircraft.index, ax
= axes[i], palette = 'viridis')

    axes[i].set_title(f"Top 5 Aircraft in '{phase}' Phase Accidents")
```

```
axes[i].set_xlabel("Number of Accidents")
axes[i].set_ylabel("Aircraft Type")
axes[i].grid(axis = 'x', linestyle = '--', alpha = 0.4)
```

*# Adjust layout*

```
plt.tight_layout()
plt.savefig('Images/Flight_phase.png', dpi=300, bbox_inches='tight') #
saves visual to image folder of my project
plt.show()
```



From the barchart above the Cessna 152 records the highest number of accidents during landing, takeoff, and cruise phases. It consistently appears as the leading aircraft involved in accidents across multiple flight stages.

## Accidents by Number of Engines

This visualization explores how accident frequency varies based on the number of engines on an aircraft. It provides insight into whether single-engine or multi-engine aircraft are more commonly involved in reported incidents

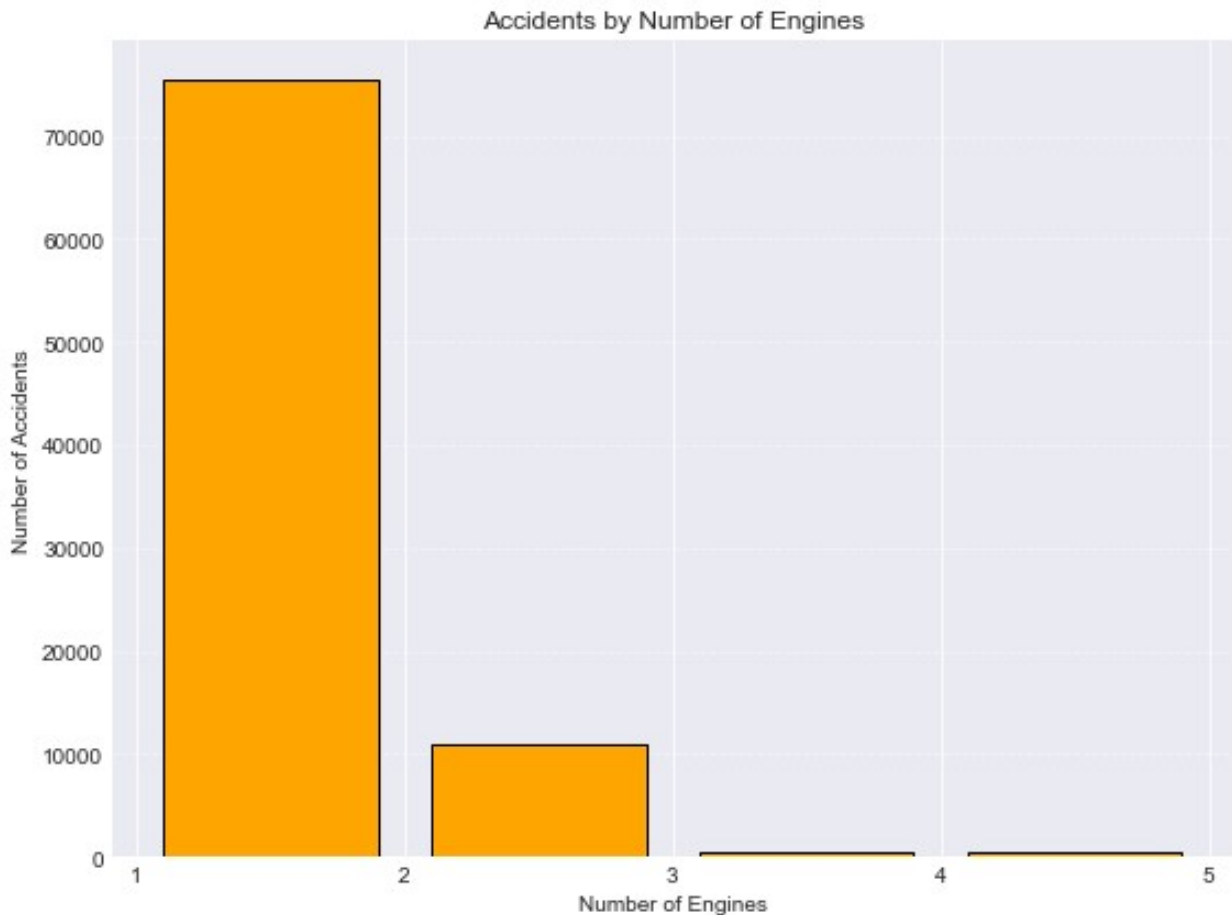
*# Plot figure*

```
plt.figure(figsize=(8, 6))
```

*# Use histogram to plot*

```
df['Number_of_engines'].plot(kind = 'hist', bins = range(1, 6), rwidth
= 0.8, color = 'orange', edgecolor = 'black')
```

```
plt.title('Accidents by Number of Engines')
plt.xlabel('Number of Engines')
plt.ylabel('Number of Accidents')
plt.grid(axis = 'y', linestyle = '--', alpha = 0.5)
plt.xticks(range(1, 6))
plt.tight_layout()
plt.savefig('Images/engines.png', dpi=300, bbox_inches='tight') #
saves visual to image folder of my project
plt.show()
```



The majority of accidents involved single-engine aircraft, accounting for over 70,000 incidents, while aircraft with two or more engines had significantly fewer accidents.

### Aircraft Accident Patterns by Purpose of Flight

This section analyzes accident patterns based on the purpose of flight, focusing on personal, instructional, business, and executive operations.

```
# Clean and standardize purpose of flight column
df['Purpose_clean'] = df['Purpose_of_flight'].str.upper().str.strip()
```

```

# Define relevant purposes for the company purpose
relevant_purposes = ['PERSONAL', 'INSTRUCTIONAL', 'BUSINESS',
'EXECUTIVE/CORPORATE']

# Filter dataset for relevant purposes
df_relevant = df[df['Purpose_clean'].isin(relevant_purposes)]

# Group by purpose and aircraft type
grouped = df_relevant.groupby(['Purpose_clean',
'Aircraft_Type']).size().reset_index(name = 'Count')


# Get top 5 aircraft for each purpose
top5_per_purpose = (grouped.groupby('Purpose_clean').apply(lambda x:
x.nlargest(5, 'Count')).reset_index(drop = True))

# Set up plot 2 rows, 2 columns
fig, axes = plt.subplots(2, 2, figsize = (14, 8))
axes = axes.flatten()

# Plot one chart per purpose
for i, purpose in enumerate(relevant_purposes):
    data = top5_per_purpose[top5_per_purpose['Purpose_clean'] ==
purpose]
    sns.barplot(x = 'Count', y = 'Aircraft_Type', data = data, ax =
axes[i], palette = 'Dark2')
    axes[i].set_title(f"Top 5 Aircraft in '{purpose.title()}'
Flights")
    axes[i].set_xlabel("Number of Accidents")
    axes[i].set_ylabel("Aircraft Type")
    axes[i].grid(axis = 'x', linestyle = '--', alpha = 0.5)

# Show
plt.tight_layout()
plt.savefig('Images/flight.png', dpi=300, bbox_inches='tight') # saves
visual to image folder of my project
plt.show()

```



The Cessna 152 and 172, 172N dominate accidents under both Instructional and Personal flight purposes. The Beech A36 and Piper PA-18 are the most involved in accidents during business flights, while the Cessna 421C and Beech 200 appear more frequently in executive or corporate flight accidents

## Business Recommendation

### 1. Accident Trends Over Time (High vs. Low-Risk Aircraft)

Recommendation: Favor the low-incident aircraft for early adoption. Approach high-incident models like the Cessna 152 and 172 with caution—only include them if comprehensive safety training and maintenance frameworks are established.

### 2. Accidents by Phase of Flight

Recommendation: The takeoff, landing, and cruise phases are the most accident-prone, with the Cessna 152 frequently involved therefore, select aircraft with proven stability and safety during these critical flight phases. Emphasize scenario-based simulation training for pilots to handle real-world challenges during these moments.

### 3. Fatalities Distribution

Recommendation: While light aircraft have higher accident frequency, Boeing 737 and 737-200 account for the highest number of fatalities per incident due to their large passenger capacity and commercial nature. Therefore, the company should Consider starting with smaller jets to build operational maturity before scaling to larger, high-capacity aircraft.



#### 4. Number of Engines

Recommendation: Single-engine aircraft dominate the accident statistics, indicating higher vulnerability during mechanical failure or adverse flight conditions. For the company entering the aviation sector, especially in commercial sector, it is advisable to prioritize multi-engine aircraft in the initial purchase.

#### 5. Flight Purpose

Cessna 152, 172, 172N dominate accident counts in Instructional and Personal flights. Beech A36 and Piper PA-18 show frequent accidents in business flights. Cessna 421C and Beechcraft 200 are commonly involved in executive or corporate. Based on the analysis, commercial aviation particularly executive and corporate flights has lower accident rates compared to personal and instructional flying. Entering the commercial sector allows the company to build trust, attract premium clients and scale operations more sustainably.

## Summary

The analysis of aircraft accident trends, fatality rates, and flight purposes suggests that entering the commercial aviation sector particularly executive and corporate operations is the most strategic and safety aligned choice. High-incident aircraft like the Cessna 152 and 172 should be approached cautiously and the company should prioritize low-incident, multi-engine aircraft with strong safety records, especially during critical phases of flight. To minimize risk and build operational maturity, I advise the company to start with smaller, safer jets, implement scenario-based pilot training, and gradually scale operations positioning itself as a reliable, safety-first aviation provider.